
Inference in Probabilistic Answer Set Programs with Imprecise Probabilities via Optimization

Damiano Azzolini¹

Fabrizio Riguzzi²

¹Department of Environmental and Prevention Sciences, University of Ferrara , Ferrara, Italy

²Department of Mathematics and Computer Science, University of Ferrara , Ferrara, Italy

Abstract

Probabilistic answer set programming has recently been extended to manage imprecise probabilities by means of credal probabilistic facts and credal annotated disjunctions. This increases the expressivity of the language but, at the same time, the cost of inference. In this paper, we cast inference in probabilistic answer set programs with credal probabilistic facts and credal annotated disjunctions as a constrained nonlinear optimization problem where the function to optimize is obtained via knowledge compilation. Empirical results on different datasets with multiple configurations show the effectiveness of our approach.

1 INTRODUCTION

Uncertainty pervades every aspect of every day’s life. Representing these situations with an expressive yet comprehensible language is crucial to understand them. Logic languages, such as Prolog [Sterling and Shapiro, 1994] and Answer Set Programming (ASP) [Brewka et al., 2011] are considered interpretable by design, since users can often encode the domain of interest with few lines of codes. In particular, ASP can compactly describe combinatorial problems. However, these languages can only represent certain data. Several semantics have been proposed to extend logic languages with constructs to represent uncertain data. One of the first was the distribution semantics (DS) [Sato, 1995] that gave the birth to the field of Probabilistic Logic Programming (PLP) and to languages such as ProbLog [De Raedt et al., 2007] and Logic Programs with Annotated Disjunctions (LPADS) [Vennekens et al., 2004]. The credal semantics (CS) [Cozman and Mauá, 2020] applies the ideas of the DS to ASP, obtaining Probabilistic Answer Set Programming (PASP).

Inference in PLP and PASP can be performed via knowl-

edge compilation [Darwiche and Marquis, 2002], where the program is converted in an alternative form which allows inference in (possibly) faster way. Recently, Mauá and Cozman [2023] extended PASP with imprecise probabilities, allowing the representation of any credal network with finitely generated credal sets, and propose a solver called dpasp based on vertex enumeration.

In this paper, we propose to solve the inference task in probabilistic answer set programs with imprecise probabilities via optimization, by extracting an equation from the result of knowledge compilation and optimizing it subject to some constraints. Empirical results on 4 different datasets with multiple configurations show that this approach is significantly faster than enumeration adopted in Mauá and Cozman [2023]. Furthermore, our solver is also able to manage annotated disjunctions with imprecise probabilities, which dpasp cannot manage.

The paper is structured as follows: Section 2 introduces background concepts involving PLP and PASP. In Section 3 we discuss how to cast inference in probabilistic answer set programs with imprecise probabilities as an optimization problem. Section 4 presents the experiments conducted to assess the performance of the developed solver. Related works are surveyed in Section 5 and Section 6 concludes the paper.

2 BACKGROUND

ProbLog [De Raedt et al., 2007] allows probabilistic facts of the form

$$\Pi :: a$$

where a is a ground atom and $\Pi \in [0, 1]$ is its probability, with the meaning that a is true (resp. false), with probability Π (resp. $1 - \Pi$). The distribution semantics (DS) [Sato, 1995] is based on the concept of *world*, a normal logic program obtained by including or not each probabilistic fact. If there are l probabilistic facts, then the number of worlds is 2^l . The

probability of a world w is computed as

$$P(w) = \prod_{a_i \in w} \Pi_i \cdot \prod_{a_i \notin w} (1 - \Pi_i) \quad (1)$$

where the $\Pi_i :: a_i$ are the probabilistic facts. The probability of an atom q , called *query*, is computed as the sum of the probabilities of the worlds where q is true. That is,

$$P(q) = \sum_{w \models q} P(w) \quad (2)$$

To clarify, consider Example 1.

Example 1 *The following probabilistic logic program has two probabilistic facts: a with probability 0.3 and b with probability 0.4.*

```
0.3 :: a.
0.4 :: b.
q :- a.
q :- b.
```

It has four worlds: w_1 where both a and b are false, with $P(w_1) = (1 - 0.3) \cdot (1 - 0.4) = 0.42$; w_2 where a is false and b is true, with $P(w_2) = (1 - 0.3) \cdot 0.4 = 0.28$; w_3 where a is true and b is false, with $P(w_3) = 0.3 \cdot (1 - 0.4) = 0.18$; and w_4 where both a and b are true, with $P(w_4) = 0.3 \cdot 0.4 = 0.12$. The probability of the query q is given by $P(q) = P(w_2) + P(w_3) + P(w_4) = 0.28 + 0.18 + 0.12 = 0.58$, since q is true in all the worlds except for w_1 .

Annotated disjunctions were introduced by Vennekens et al. [2004] with the syntax $h_1 : \Pi_1; \dots; h_m : \Pi_m :- b_1, \dots, b_n$, where $\sum_i \Pi_i = 1$. The meaning is that, when the conjunction of the literals b_j in the body is true, one of the head atoms h_i is true with the corresponding probability Π_i . We consider the notation

$$\Pi_1 :: h_1; \dots; \Pi_m :: h_m :- b_1, \dots, b_n$$

for uniformity with ProbLog. Annotated disjunctions can be converted into probabilistic facts [De Raedt et al., 2008] as follows: for each annotated disjunction $\Pi_1 :: h_1; \dots; \Pi_m :: h_m :- B$ (with B the body) with m heads we add $m - 1$ probabilistic facts and m rules:

$$\begin{aligned} \pi_1 &:: f_1. \\ \dots \\ \pi_{m-1} &:: f_{m-1}. \\ h_1 &:- B, f_1. \\ h_2 &:- B, \text{not } f_1, f_2. \\ \dots \\ h_m &:- B, \text{not } f_1, \dots, \text{not } f_{m-1}. \end{aligned} \quad (3)$$

where $\pi_1 = \Pi_1$ and $\pi_i = \Pi_i / \sum_{j=1}^{i-1} (1 - \pi_j)$ for $i > 0$. For example, the probabilistic facts obtained by converting

the annotated disjunction $0.2 :: a; 0.3 :: b; 0.5 :: c$ have probability 0.2 and $0.3/(1 - 0.2) = 0.375$.

The credal semantics (CS) assigns a meaning to probabilistic answer set programs [Cozman and Mauá, 2020], i.e., answer set programs [Brewka et al., 2011] extended with ProbLog probabilistic facts. In this setting, a world may have zero or more stable models (also called answer sets). A stable model is a minimal model under set inclusion of the reduct of an answer set program P , where the reduct of P w.r.t. an interpretation I is obtained by removing from P all the rules that have the body false in I . An interpretation is called model if it satisfies all the groundings of the rules of P . The CS requires that every world has at least one answer set. The probability of a query q under the CS is described by a probability range, defined by a lower ($\underline{P}(q)$) and an upper ($\overline{P}(q)$) bound. A world w contributes to both the lower and upper bound if the query is true in *every* answer set of w . Conversely, a world w contributes to only the upper bound if the query is true in *at least one* answer set of w . This can be expressed in formulas as

$$\underline{P}(q) = \sum_{w_i | \forall m \in AS(w_i), m \models q} P(w_i) \quad (4)$$

$$\overline{P}(q) = \sum_{w_i | \exists m \in AS(w_i), m \models q} P(w_i) \quad (5)$$

Example 2 *The following probabilistic answer set program is a variation of the probabilistic logic program shown in Example 1.*

```
0.3 :: a.
0.4 :: b.
q :- a.
q ; r :- b.
```

Consider, as in Example 1, the query q . w_2 (where a is false and b is true) has 2 answer sets, $\{q, b\}$ and $\{r, b\}$, and it only contributes to the upper probability for q , since q is true only in the first. w_3 and w_4 have a unique answer set each ($\{a, q\}$ and $\{a, b, q\}$, respectively) where q is true, so they contribute to both lower and upper probability. w_1 has a unique answer set with no atoms, so it does not contribute. Overall, $P(q) = [\underline{P}(q), \overline{P}(q)] = [P(w_3) + P(w_4), P(w_2) + P(w_3) + P(w_4)] = [0.3, 0.58]$.

Inference in probabilistic answer set programs can be performed via Second Level Algebraic Model Counting (2AMC) [Kiesel et al., 2022]. 2AMC is an abstract framework that comprises several well-known problems, such as inference, decision theoretic inference, and MAP inference. The task considers two commutative semirings [Gondran and Minoux, 2008] $R^i = \{D^i, \oplus^i, \otimes^i, n_{\oplus^i}, n_{\otimes^i}\}$ and $R^o = \{D^o, \oplus^o, \otimes^o, n_{\oplus^o}, n_{\otimes^o}\}$, a propositional theory T whose variables are split into two disjoint sets, V_o and V_i ,

two weight functions w_i and w_o , and a transformation function f , and requires solving:

$$2AMC(T) = \bigoplus_{I_o \in \mu(V_o)}^o \bigotimes_{a \in I_o}^o w_o(a) \otimes^o f\left(\bigoplus_{I_i \in \varphi(T|I_o)}^i \bigotimes_{b \in I_i}^i w_i(b)\right) \quad (6)$$

where $\mu(V_o)$ is the set of possible assignments to the variables in V_o and $\varphi(T|I_o)$ is the set of possible assignments to the variables in T that satisfy I_o . Said otherwise, for every assignment of the variables V_o , we need to solve an inner AMC [Kimmig et al., 2017] task on the variables V_i . The result of the inner AMC is converted through the transformation function f into a value coherent with the ones of the variables in the outer semiring and we need to solve a second AMC task, this time by considering V_o . Kiesel et al. [2022] also introduced a tool called *aspmc* [Eiter et al., 2021] to solve 2AMC problems, based on a conversion guided by the treewidth of the program into a compact form via knowledge compilation [Darwiche and Marquis, 2002]. The target of the compilation is negation normal form (NNF), a rooted directed acyclic graph where internal nodes are associated with conjunction (AND-nodes) or disjunctions (OR-nodes) and leaves are associated with literals, true, or false. Furthermore, the obtained NNF has the properties of decomposability, determinism, smoothness, and X-first modulo definability. In particular, the last property imposes an hierarchical structure on the appearance of the variables, allowing to consider them according to the specified semiring. This tree can be traversed bottom up to solve the 2AMC task.

Azzolini and Riguzzi [2023] introduced *aspcs*, a solver built on top of *aspmc*, to perform inference in probabilistic answer set programs. To do so, they instantiate Equation 6 by considering as inner semiring $\mathcal{R}_{in} = (\mathbb{N}^2, +, \cdot, (0, 0), (1, 1))$ with w_i mapping *not* q (assuming that we are interested in computing the probability of the query q) to $(0, 1)$ and all other literals to $(1, 1)$, where operations are performed component-wise. In this inner semiring, we have a fixed world. Practically, \mathcal{R}_{in} counts the number of answer sets where the query is true and the total number of answer sets. The transformation function is $f(n_1, n_2)$ returning a pair (v_{lp}, v_{up}) where $v_{lp} = 1$ if $n_1 = n_2$, 0 otherwise, and $v_{up} = 1$ if $n_1 > 0$, 0 otherwise. This function checks the bounds to which the considered world contributes. The outer semiring is an extension of the probability semiring [Kimmig et al., 2017] to two dimensions, i.e., $\mathcal{R}_{out} = ([0, 1]^2, +, \cdot, (0, 0), (1, 1))$ (the operations are still considered component-wise), with w_o associating (p, p) and $(1 - p, 1 - p)$ to a and *not* a , respectively, for every probabilistic fact $p :: a$ and $(1, 1)$ to all the remaining literals. In other words, in \mathcal{R}_{out} , we multiply the probabilities of the probabilistic facts and sum them to obtain the probability of q .

Mauá and Cozman [2023] introduced probabilistic facts

with an imprecise probability, that in this paper we call *credal facts*. These are expressed with the syntax $[\alpha, \beta] :: a$ with the meaning that a (a ground atom) has a probability ranging between α and β , with $\alpha, \beta \in \mathbb{R}$, $0 \leq \alpha \leq \beta \leq 1$. Mauá and Cozman [2023] also discussed annotated disjunctions with imprecise probabilities, that we call *credal annotated disjunctions*, for uniformity with credal facts. Their syntax is

$$[\alpha_1, \beta_1] :: h_1; \dots; [\alpha_m, \beta_m] :: h_m :- b_1, \dots, b_n \quad (7)$$

where each h_i is an atom and each b_j is a literal. The meaning is that, when the conjunction of the literals in the body is true, one of the head atoms is true with the corresponding probability range. To ensure well definedness, the reachability property must be met, i.e., $\alpha_i + \sum_{j \neq i} \beta_j \geq 1$ and $\beta_i + \sum_{j \neq i} \alpha_j \leq 1$ for all i .

3 INFERENCE WITH CREDAL FACTS AND CREDAL ANNOTATED DISJUNCTIONS

In this section, we first discuss how to cast inference in programs with credal facts only (and possibly probabilistic facts) as an optimization problem. Then, we extend that approach to also manage credal annotated disjunctions.

3.1 INFERENCE WITH CREDAL FACTS

Inference in programs with credal facts and without credal annotated disjunctions can be cast as inference in probabilistic answer set programs with only probabilistic facts, by converting (as described by [Mauá and Cozman, 2023]) each credal fact $[\alpha, \beta] :: a$ into an annotated disjunction $\alpha :: a1; 1 - \beta :: a2; \beta - \alpha :: a3$ and three rules $a :- a1$, $na :- a2$, and $a; na :- a3$. To clarify, consider Example 3.

Example 3 *The following probabilistic answer set program is a variation of one shown in Example 2. Now, a and b are credal facts.*

```
[0.3, 0.4] :: a.
[0.4, 0.9] :: b.
q :- a.
q ; r :- b.
```

This program is converted into

```
0.3 :: a1 ; 0.6 :: a2 ; 0.1 :: a3.
a :- a1.
na :- a2.
a ; na :- a3.
0.4 :: b1 ; 0.1 :: b2 ; 0.5 :: b3.
b :- b1.
nb :- b2.
```

```

b ; nb :- b3.
q :- a.
q ; r :- b.

```

The probability of the query q is $[0.3, 0.94]$.

However, this conversion greatly increases the number of probabilistic facts and rules, possibly slowing down the inference process. We propose an alternative approach based on 2AMC. We consider as inner semiring and transformation function the same of Azzolini and Riguzzi [2023] and described in Section 2. For the outer semiring \mathcal{R}_{out} we propose an extension to two dimensions of the sensitivity semiring [Kimmig et al., 2017]: $\mathcal{R}_{out} = (\mathbb{R}[X], +, -, (0, 0), (1, 1))$, with w_o associating (p_i, p_i) and $(1 - p_i, 1 - p_i)$ to a_i and *not* a_i , respectively, for every probabilistic fact $p_i :: a_i$; (π_i, π_i) and $(1 - \pi_i, 1 - \pi_i)$ to v_i and *not* v_i , respectively, for every credal fact $[\alpha, \beta] :: v_i$, where the π_i are symbolic variables representing the probability of the credal facts. The remaining literals are assigned $(1, 1)$. In this way, by solving the 2AMC, we obtain two non-linear formulas with at most k variables, where k is the number of credal facts, composed by summations and products only (except for trivial cases). Let us call them $f_{lp}(X)$ and $f_{up}(X)$, where $X = \{\pi_1, \dots, \pi_k\}$. Their general form is $f_{lp/up}(X) = \sum_w \prod_{a_i \in w} \pi_i \prod_{a_j \notin w} (1 - \pi_j)$. If we replace each variable with a value, we get the lower and upper probability for the query when the facts have the chosen values. Now, the lower probability for q can be obtained by minimizing $f_{lp}(X)$ where the variables π_1, \dots, π_k have the range described by the credal fact they represent. That is, if π_1 is associated to the credal fact $[0.2, 0.6] :: a_i$, π_1 must be in the range $[0.2, 0.6]$. Similarly, the upper probability for q can be obtained by maximizing $f_{up}(X)$ (or, equivalently, minimizing $-f_{up}(X)$) where the bounds of the variables are the same as before. In other words, we cast inference in programs with credal facts as a nonlinear optimization problem. That is, $\underline{P}(q)$ can be obtained by solving

$$\begin{aligned} & \text{minimize } f_{lp}(X) \\ & \text{s.t. } \pi_i \in [l_i, u_i], \forall i \in \{1, \dots, k\} \end{aligned} \quad (8)$$

while $\bar{P}(q)$ by solving

$$\begin{aligned} & \text{maximize } f_{up}(X) \\ & \text{s.t. } \pi_i \in [l_i, u_i], \forall i \in \{1, \dots, k\} \end{aligned} \quad (9)$$

In this way, we avoid the explosion of the size of the program due to the introduction of auxiliary annotated disjunctions.

Example 4 Consider the program with credal facts shown in Example 3. It is converted into

```

pa :: a.
pb :: b.
q :- a.
q ; r :- b.

```

By traversing the NNF for the query q we extract two equations: $f_{lp}(pa) = pa$ and $f_{up}(pa, pb) = pa - pb \cdot (pa - 1)$, the former for the lower probability and the latter for the upper probability. To compute $\underline{P}(q)$, we need to minimize $f_{lp}(pa)$ with $pa \in [0.3, 0.4]$. Clearly, the minimum value of $f_{lp}(pa)$ is 0.3. Consider now the computation of $\bar{P}(q)$. We need to maximize $f_{up}(pa, pb)$, or equivalently, minimize $-f_{up}(pa, pb)$ with $pa \in [0.3, 0.4]$ and $pb \in [0.4, 0.9]$. In this case, the maximum value of $f_{up}(pa, pb)$ is 0.94. Thus, $[\underline{P}(q), \bar{P}(q)] = [0.3, 0.94]$.

3.2 INFERENCE WITH CREDAL ANNOTATED DISJUNCTIONS

When credal annotated disjunction are present in a program, the conversion between annotated disjunctions and probabilistic facts (see Section 2) does not preserve the equivalence, as discussed in [Mauá and Cozman, 2023]. Still in [Mauá and Cozman, 2023], to perform inference, the authors propose to convert each credal annotated disjunction with k disjunctions in the head into a set of k credal facts with vacuous intervals (i.e., $\alpha = 0$ and $\beta = 1$), a set of deterministic rules, and a set of annotated disjunctions representing all the vertices of the induced credal set. The set of annotated disjunctions (as well as the set of deterministic rules) obtained with this conversion may be exponential in the number of credal facts, making inference very expensive. If we consider the example discussed in [Mauá and Cozman, 2023], the credal annotated disjunction $D = [0.1, 0.3] :: \text{red}; [0.2, 0.4] :: \text{green}; [0.4, 0.6] :: \text{blue}$, is converted into three credal facts with vacuous intervals, 10 deterministic rules, and 6 annotated disjunctions with 3 atom each in the head.

Here we follow another path: we convert each credal annotated disjunction into a set of probabilistic facts and rules, as explained in Section 2. However, we do not compute the probabilities of the obtained probabilistic facts but we leave them symbolic. More precisely, for a credal annotated disjunction of the form of Equation 7, we get the rules of Equation 3. That is, from D discussed few lines above we obtain: $\pi_1 :: f_1$, $\pi_2 :: f_2$, $\text{red} :- f_1$, $\text{green} :- \text{not } f_1, f_2$, and $\text{blue} :- \text{not } f_1, \text{not } f_2$. Note again that the probability of f_1 and f_2 is kept symbolic. To compute the probability of a query q , we extract, as described in Section 3.1, two symbolic equations, one for the lower, $f_{lp}(X)$, and one for the upper probability, $f_{up}(X)$. As before, we treat the inference process as an optimization problem: $\underline{P}(q)$ can be computed by minimizing $f_{lp}(X)$ while $\bar{P}(q)$ by maximizing $f_{up}(X)$. Both problems, however, require imposing an additional set of nonlinear constraints that mimics the probability conversion between annotated disjunctions and probabilistic facts (Section 2). With n_{ad} credal annotated disjunctions of the form of Equation 7, the constraints can be compactly

expressed as

$$\begin{aligned}
& \text{minimize} && f(X) \\
& \text{s.t.} && \pi_i^l \cdot \prod_{j < i} (1 - \pi_j^l) - \alpha_i^l \geq 0, \\
& && \beta_i^l - \pi_i^l \cdot \prod_{j < i} (1 - \pi_j^l) \geq 0, \\
& && \forall l \in \{1, \dots, n_{ad}\}, \forall i \in \{1, \dots, m_l\}
\end{aligned}$$

assuming $\pi_i^{m_i} = 1$, where π_i^k is the probability associated to the i -th probabilistic fact related to the i -th head of the k -th annotated disjunction. Overall, for a credal annotated disjunction with m heads we add $2 \cdot m$ constraints. So, the problem of inference can be cast as two nonlinear optimization problems with nonlinear constraints. For example, with a credal annotated disjunction of the form $[\alpha_1, \beta_1] :: h_1; [\alpha_2, \beta_2] :: h_2; [\alpha_3, \beta_3] :: h_3$ call π_1 and π_2 the probabilities associated with the probabilistic facts obtained via the conversion discussed in Section 2, we have $\pi_1 - \alpha_1 \geq 0$, $\beta_1 - \pi_1 \geq 0$, $(1 - \pi_1) \cdot \pi_2 - \alpha_2 \geq 0$, $\beta_2 - (1 - \pi_1) \cdot \pi_2 \geq 0$, $(1 - \pi_1) \cdot (1 - \pi_2) - \alpha_3 \geq 0$, and $\beta_3 - (1 - \pi_1) \cdot (1 - \pi_2) \geq 0$. More concretely, with the credal annotated disjunction D , we obtain the set of constraints: $c_1 - 0.1 \geq 0$, $0.3 - \pi_1 \geq 0$, $(1 - \pi_1) \cdot \pi_2 - 0.2 \geq 0$, $0.4 - (1 - \pi_1) \cdot \pi_2 \geq 0$, $(1 - \pi_1) \cdot (1 - \pi_2) - 0.4 \geq 0$, and $0.6 - (1 - \pi_1) \cdot (1 - \pi_2) \geq 0$. Note that this approach can be straightforwardly extended to support *parameterized* annotated disjunctions, i.e., annotated disjunctions where the probabilities of the heads are specified via a set of constraints, for example $A :: a; B :: b; C :: c :- A < B, B < C, C > 0.1$, where A, B and C are the probabilities to be determined. In our framework, this consists in adding the set of constraints specified in the body to the optimization problem.

4 EXPERIMENTS

We implemented the proposed algorithm in Python on top of the aspcs solver [Azzolini and Riguzzi, 2023] and leveraged the SciPy library [Virtanen et al., 2020] to solve the optimization problems and SymPy [Meurer et al., 2017] to simplify the equation extracted from the NNF¹. We tested both the COBYLA [Powell, 1994] and SLSQP [Kraft, 1994] algorithms for nonlinear constrained optimization, without changing the default parameters. We compare it against dpasp [Mauá and Cozman, 2023], which adopts vertex enumeration and only supports credal facts (not credal annotated disjunctions). We run the evaluation on a machine running at 2.40 GHz with an execution time limit of 8 hours and 16 GB of RAM.

We set up two different set of experiments: in a first round, call it E_1 , we consider four datasets with probabilistic and

credal facts and compare aspcs against dpasp; in a second round, call it E_2 , we consider a variation of the same four datasets but with credal annotated disjunctions instead of credal facts. dpasp does not support these, so we only report the results for aspcs. All the instances for all the datasets have at least one answer set per world, as required by the credal semantics.

Let us start by describing the four datasets for E_1 , adapted from [Azzolini and Riguzzi, 2023]. The programs for aspcs and dpasp only differ in the negation symbol: $\backslash +$ for the former and *not* for the latter. The following snippets show the aspcs version. The first dataset, *coloring*, encodes a graph coloring scenario, a well-known problem that can be easily modeled in ASP. Here, some nodes can be associated with three distinct colors, namely red, green, and blue, and others have a fixed color. Nodes are connected by credal probabilistic facts *edge/2* and connected nodes must have different colors. All the instances have the same following rules:

```

red(X) :- node(X), \+ green(X),
        \+ blue(X).
green(X) :- node(X), \+ red(X),
        \+ blue(X).
blue(X) :- node(X), \+ red(X),
        \+ green(X).
e(X,Y) :- edge(X,Y).
e(Y,X) :- edge(Y,X).
:- e(X,Y), red(X), red(Y).
:- e(X,Y), green(X), green(Y).
:- e(X,Y), blue(X), blue(Y).

```

but they differ in the number of nodes and number of edges (this value denotes the size of the instance). Each instance has an additional set of rules $qr :- blue(i), \forall i \in \{1, \dots, n\}$, where n is the number of nodes. The query is *qr*.

The *smoke* dataset, introduced in Totis et al. [2023], encodes a network of people where some of them are smokers while other smoke only due to the influence of their friends. All the programs have the following rules:

```

asthma_r(X) :- smokes(X), asthma_f(X).
asthma(X) :- asthma_f(X).
asthma(X) :- asthma_r(X).
asthma_and_stress(X) :-
    stress(X), stress_f(X), asthma(X).
smokes(X) :- infl(Y,X), smokes(Y).
smokes(X) :- asthma_and_stress(X),
        \+ no_smokes(X).
no_smokes(X) :- asthma_and_stress(X),
        \+ smokes(X).

```

where *asthma_f/1*, *stress/1*, *stress_f/1*, and *asthma_f/1* are probabilistic with a sharp probabil-

¹Source code and datasets available at: <https://github.com/damianoazzolini/aspmc>.

ity value associated (and there is one of them for each people involved) while the *infl/2* facts are credal probabilistic facts. The instances have an increasing number of people, and thus an increasing number of probabilistic and credal facts. The query is *smokes(1)*.

The dataset *irl* contains a set of instances with three clauses with an increasing number of credal probabilistic facts a_i in the body: $qr :- \bigwedge_{i < k, i \text{ even}} a_i$, $qr :- \bigwedge_{i < k, i \text{ odd}} a_i, \setminus + nqr$, and $nqr :- \bigwedge_{i < k, i \text{ odd}} a_i, \setminus + qr$, where k denotes the size of the instance. The instance of size 5 is

```
qr:- a0, a2, a4.
qr:- a1, a3, \+ nqr.
nqr:- a1, a3, \+ qr.
```

where the $a_i, i \in \{0, \dots, 4\}$, are credal facts. The query is *qr*.

Each instance k of the fourth dataset, *irn*, contains a set of $qr :- a_i$ rules $\forall i \in \{0, \dots, k-1\}$ even, $qr :- a_i, \setminus + nqr$ $\forall i \in \{0, \dots, k-1\}$ odd, and $nqr :- a_i, \setminus + qr$ $\forall i \in \{0, \dots, k-1\}$ odd. For example, the instance of size 5 is

```
qr:- a0. qr:- a2. qr:- a4.
qr :- a1, \+ nqr. nqr :- a1, \+ qr.
qr :- a3, \+ nqr. nqr :- a3, \+ qr.
```

The query is *qr*.

For all the datasets we consider two variations, called *loose* and *strict*: for the former, all the credal facts are associated with the probability range [0.05,0.95]; for the latter, the range is [0.45,0.55]. This is because we want to evaluate the optimization process with different types of constraints.

Since the equations extracted from the NNF usually involve many products and summations and this may possibly slow down the optimization process, in a preliminary test on the *coloring* dataset we compared the execution time of *aspcs* with and without simplifying them by using *SymPy*. In particular, we are interested in assessing whether the extra execution time spent to simplify the equations effectively speeds up the overall execution time. Table 1 shows the total execution time of running *aspcs* with the COBYLA algorithm on the *coloring* dataset in both strict and loose version with and without simplification: for the largest solvable instance (size 16) the execution time with simplification is 6 times less than the one without. This proves that the simplification is a crucial step in the pipeline. Note that the simplification is performed twice, once for the lower and once for the upper probability. So, in the remaining experiments, we always consider *aspcs* with simplification.

To investigate even further the impact of the different steps involved, Table 2 reports the NNF computation time, equations simplification time, and optimizations time for the COBYLA algorithm on the loose configuration of the *coloring* dataset. For bigger instances, the simplification time

Table 1: Impact of the simplification process on the probability computation with the COBYLA algorithm on the *coloring* dataset. The columns represent, respectively, the size of the instance (*size*), the execution time (seconds) for the strict version with simplification (*s. w.*), the execution time (seconds) for the strict version without simplification (*s. wo.*), the execution time (seconds) for the loose version with simplification (*l. w.*), and the execution time (seconds) for the strict version without simplification (*l. wo.*).

<i>size</i>	<i>s. w.</i>	<i>s. wo.</i>	<i>l. w.</i>	<i>l. wo.</i>
8	4.428	5.727	5.39	6.612
9	5.094	6.981	5.071	7.436
10	7.536	13.284	8.014	14.505
11	13.406	26.228	12.836	27.313
12	33.116	47.158	32.947	47.298
13	46.987	103.322	46.778	126.434
14	46.995	164.814	46.747	172.029
15	84.338	308.938	84.876	362.038
16	104.057	627.793	104.631	675.879

Table 2: Impact in seconds and percentage on the total execution time of NNF construction (*NNF*), simplification (*simpl*), and optimization (*opt*) on the probability computation with the COBYLA algorithm on the loose configuration of the *coloring* dataset. The percentages do not sum to 100 due to other internal operations.

<i>size</i>	<i>NNF(s)</i>	<i>simpl.(s)</i>	<i>opt.(s)</i>
8	2.192 (40.66%)	1.449 (26.88%)	0.058 (1.07%)
9	3.250 (64.08%)	0.546 (10.76%)	0.064 (1.26%)
10	4.399 (54.89%)	2.008 (25.05%)	0.134 (1.67%)
11	9.252 (72.07%)	2.257 (17.58%)	0.095 (0.74%)
12	29.68 (90.08%)	1.859 (5.642%)	0.076 (0.23%)
13	35.04 (74.90%)	10.27 (21.95%)	0.162 (0.34%)
14	35.88 (76.75%)	9.303 (19.90%)	0.212 (0.45%)
15	43.17 (50.86%)	39.97 (47.09%)	0.226 (0.26%)
16	46.31 (44.26%)	56.56 (54.05%)	0.321 (0.30%)

requires more than half of the total execution time (again, the reported value is the sum of two simplifications, one for the equation for the lower probability and one for the equation for the upper probability). However, this makes solving the optimizations (also here the time represents optimization applied on both equations) very quick (less than a second), compared to the hundreds of seconds required in the case the equation is not simplified (see Table 1). We report these detailed values only for the *coloring* dataset, due to space restrictions. However, these considerations apply to all the datasets.

We are now ready to discuss the results for E_1 . Figure 1 shows the results on the *coloring* dataset for all the configurations. *aspcs* is faster than *dpasp* and COBYLA and SLSQP for both loose and strict configurations have similar execution times. The instance of size 16 was the largest

solvable one. For *dpasp*, as expected, the loose and strict configurations require the same time to query. The figure also reports the execution times for the converted instances, where credal facts were converted into annotated disjunctions. Also in this case *aspcs* outperforms *dpasp*, and, for both, inference in the program with credal facts is substantially faster than inference in the converted program. This is probably due to the large number of probabilities introduced by the annotated disjunctions. In this case, the largest solvable instance had size 13 for both. For *dpasp* we only plot up to size 11 since 12 and 13 required respectively 3340 and 25083 seconds and would have made all the other curves in the figure unreadable.

Figure 2 shows the results on the *smoke* dataset. As for the *coloring* dataset, *aspcs* is faster than *dpasp*. For *dpasp*, the loose configuration of instance 6 is approximately 500 seconds faster than the strict. *aspcs* requires approximately 60 seconds to run the converted version. *dpasp* was also able to solve it in 26861 seconds (that we do not plot). Interestingly, the converted version for the strict SLSQP configuration is slightly faster than the non-converted one: this may be due to knowledge compilation that can find a good NNF structure for the program.

The results for *irl* are shown in Figure 3. Here, *aspcs* can solve the instance of size 20 with both COBYLA and SLSQP in both configurations in less than 10 seconds, while *dpasp* requires more than 1500 seconds. For the converted version, the largest solvable instance size was 15: also in this case *aspcs* shows better performances than *dpasp*. This proves that knowledge compilation can find a good representation of the program, since, for example, the equation for the upper probability is composed by the product of all the probabilities of the credal facts.

Lastly, Figure 4 shows the results for *irn*: it confirms the trend of the three previous datasets, where *aspcs* is faster than *dpasp* and the converted version is slower than the version with credal probabilistic facts.

For E_2 , we modify the four datasets to include credal annotated disjunctions. For the *coloring* dataset, we add a credal annotated disjunction $[0.1, 0.3] :: c0(X); [0.2, 0.4] :: c1(X); [0.4, 0.6] :: c2(X)$ for every node X and modify the constraint $:- e(X, Y), blue(X), blue(Y)$ into $:- e(X, Y), blue(X), blue(Y), c2(X), c2(Y)$. Furthermore, edges are now considered probabilistic facts. The remaining rules are the same. For *smoke*, as for *coloring*, we add a credal annotated disjunction $[0.1, 0.3] :: c0(X); [0.2, 0.4] :: c1(X); [0.4, 0.6] :: c2(X)$ and replace the rule $smokes(X) :- infl(Y, X), smokes(Y)$ with $smokes(X) :- infl(Y, X), smokes(Y), c0(X), c1(Y)$ and consider the *infl* facts probabilistic with a sharp probability value. Lastly, for *irl* and *irn* we replace each credal probabilistic fact a_i with: i) $[0.1, 0.3] :: a_i; [0.2, 0.4] :: a_{i1}(X); [0.4, 0.6] :: a_{i2}(X)$ if $i \% 3 = 0$; ii) $[0.1, 0.3] ::$




Figure 1: Execution times for the *coloring* datasets for *aspcs* (solid lines) and *dpasp* (dash-dotted lines). Dotted lines are the results obtained in programs where the credal facts are converted into annotated disjunctions.




Figure 2: Execution times for the *smoke* datasets for *aspcs* (solid lines) and *dpasp* (dash-dotted lines). Dotted lines represent the results obtained in the programs where the credal facts were converted into annotated disjunctions.

$a_{i1}; [0.2, 0.4] :: a_i(X); [0.4, 0.6] :: a_{i2}(X)$ if $i \% 3 = 1$; or ii) $[0.1, 0.3] :: a_{i1}; [0.2, 0.4] :: a_{i2}(X); [0.4, 0.6] :: a_i(X)$ if $i \% 3 = 2$. The rules are the same as for E_1 .

Table 3 shows the results. The algorithm can solve fewer instances than the version with only credal facts: this is due to the increasing number of probabilistic facts that need to be considered, whose number is indicated in the fourth column. Furthermore, the COBYLA algorithm is always, except for one case, faster than SLSQP.

5 RELATED WORK

Apart from the credal semantics, there exist various alternatives to represent uncertainty with answer set programs, such as LP^{MLN} [Lee and Wang, 2016], P-log [Baral et al., 2009], and diff SAT [Nickles, 2018], but imprecise probabil-




Figure 3: Execution times for the *irl* datasets for aspcs (solid lines) and dpasp (dash-dotted lines). Dotted lines represent the results obtained in the programs where the credal facts were converted into probabilistic facts.




Figure 4: Execution times for the *irm* datasets for aspcs (solid lines) and dpasp (dash-dotted lines). Dotted lines represent the results obtained in the programs where the credal facts were converted into probabilistic facts.

ities have not been proposed for these. Also the approaches of Totis et al. [2023] and Rocha and Cozman [2022] are probabilistic extensions of ASP, both targeting argumentation, but they still do not consider probabilistic fact with imprecise probabilities. As previously discussed, Mauá and Cozman [2023] introduced credal probabilistic facts and credal annotated disjunctions and developed a solver based on vertex enumeration, that, however, does not yet support credal annotated disjunctions. Azzolini and Riguzzi [2021] integrated constrained optimization within probabilistic logical inference, but they considered PLP, not PASP, and a different optimization framework. Constraints are a standard component of ASP solvers, but they are limited to integers (i.e., floating points are usually not supported, even if some alternatives are under development [Pacenza and Zangari, 2023]) and probabilities are not managed. Some works on

Table 3: Execution times in seconds for the experiments with credal annotated disjunctions. The last two columns contain the number of probabilistic facts and the number of rules in each instance, respectively. Lowest execution times are in bold.

<i>size</i>	<i>cobyala</i> (s)	<i>slsqp</i> (s)	<i>#pf</i>	<i>#rules</i>
<i>coloring</i>				
8	45.135	55.151	20	14
9	104.984	105.839	23	15
10	503.691	504.313	24	15
<i>irl</i>				
5	5.084	12.438	10	18
10	24.915	56.726	20	33
15	755.598	877.51	30	48
20	12132.332	11956.459	40	63
<i>irn</i>				
10	264.298	284.184	20	45
11	1339.725	2184.979	22	49
12	4525.672	4837.060	24	54
13	4752.085	5994.453	26	58
14	14627.982	15384.466	28	63
<i>smoke</i>				
1	9.67	11.45	14	8
2	23.025	24.024	18	8
3	168.59	172.355	22	8
4	408.532	408.542	25	8

this line can be found in [Arias et al., 2018, Lierler, 2023].

Knowledge compilation is a well-known technique in probabilistic (logical) settings, used to perform probabilistic inference [De Raedt et al., 2007, Dries et al., 2015] also in programs with both discrete and continuous random variables [Zuidberg Dos Martires et al., 2019], and to solve decision theoretic problems [Van den Broeck et al., 2010], allowing a compact representation of the program at hand.

6 CONCLUSION

In this paper, we discussed how to perform inference in probabilistic answer set programs with imprecise probabilities via optimization. Our pipeline consists of four steps: construction of a NNF representation of the query, extraction of two equations from the NNF, one for the lower and one for the upper probability, simplification of the equations, and solution of two constrained nonlinear optimization problems (minimization for the lower probability and maximization for the upper probability). Empirical results show that our approach is significantly faster than an already existing solver. Future works involve the study of related inference tasks, such MAP inference with imprecise probabilities.

Acknowledgements

This work has been partially supported by the Spoke 1 “FutureHPC & BigData” of the Italian Research Center on High-Performance Computing, Big Data and Quantum Computing (ICSC) funded by MUR Missione 4 - Next Generation EU (NGEU), and by TAILOR, a project funded by EU Horizon 2020 research and innovation programme under GA No. 952215. Both authors are members of the Gruppo Nazionale Calcolo Scientifico – Istituto Nazionale di Alta Matematica (GNCS-INdAM).

References

- Joaquín Arias, Manuel Carro, Elmer Salazar, Kyle Marple, and Gopal Gupta. Constraint answer set programming without grounding. *Theory and Practice of Logic Programming*, 18(3-4):337–354, 2018. doi: 10.1017/S1471068418000285.
- Damiano Azzolini and Fabrizio Riguzzi. Optimizing probabilities in probabilistic logic programs. *Theory and Practice of Logic Programming*, 21(5):543–556, 2021. doi: 10.1017/S1471068421000260.
- Damiano Azzolini and Fabrizio Riguzzi. Inference in probabilistic answer set programming under the credal semantics. In Roberto Basili, Domenico Lembo, Carla Limongelli, and Andrea Orlandini, editors, *AIxIA 2023 - Advances in Artificial Intelligence*, volume 14318 of *Lecture Notes in Artificial Intelligence*, pages 367–380, Heidelberg, Germany, 2023. Springer. doi: 10.1007/978-3-031-47546-7_25.
- Chitta Baral, Michael Gelfond, and Nelson Rushton. Probabilistic reasoning with answer sets. *Theory and Practice of Logic Programming*, 9(1):57–144, 2009. doi: 10.1017/S1471068408003645.
- Gerhard Brewka, Thomas Eiter, and Mirosław Truszczyński. Answer set programming at a glance. *Communications of the ACM*, 54(12):92–103, december 2011. ISSN 0001-0782. doi: 10.1145/2043174.2043195.
- Fábio Gagliardi Cozman and Denis Deratani Mauá. The joy of probabilistic answer set programming: Semantics, complexity, expressivity, inference. *International Journal of Approximate Reasoning*, 125:218–239, 2020. ISSN 0888-613X. doi: 10.1016/j.ijar.2020.07.004.
- Adnan Darwiche and Pierre Marquis. A knowledge compilation map. *Journal of Artificial Intelligence Research*, 17:229–264, 2002. doi: 10.1613/jair.989.
- Luc De Raedt, Angelika Kimmig, and Hannu Toivonen. ProbLog: A probabilistic Prolog and its application in link discovery. In Manuela M. Veloso, editor, *IJCAI 2007*, volume 7, pages 2462–2467. AAAI Press, 2007.
- Luc De Raedt, Bart Demoen, Daan Fierens, Bernd Gutmann, Gerda Janssens, Angelika Kimmig, Niels Landwehr, Theofrastos Mantadelis, Wannes Meert, Ricardo Rocha, Vítor Costa, Ingo Thon, and Joost Vennekens. Towards digesting the alphabet-soup of statistical relational learning. In *NIPS 2008 Workshop on Probabilistic Programming*, 2008.
- Anton Dries, Angelika Kimmig, Wannes Meert, Joris Renkens, Guy Van den Broeck, Jonas Vlasselaer, and Luc De Raedt. ProbLog2: Probabilistic logic programming. In *European Conference on Machine Learning and Principles and Practice of Knowledge Discovery in Databases (ECMLPKDD 2015)*, volume 9286 of *LNCS*, pages 312–315. Springer, 2015. doi: 10.1007/978-3-319-23461-8_37.
- Thomas Eiter, Markus Hecher, and Rafael Kiesel. Treewidth-aware cycle breaking for algebraic answer set counting. In Meghyn Bienvenu, Gerhard Lakemeyer, and Esra Erdem, editors, *Proceedings of the 18th International Conference on Principles of Knowledge Representation and Reasoning, KR 2021*, pages 269–279, 2021. doi: 10.24963/kr.2021/26.
- Michel Gondran and Michel Minoux. *Graphs, Dioids and Semirings: New Models and Algorithms (Operations Research/Computer Science Interfaces Series)*. Springer Publishing Company, Incorporated, 1 edition, 2008. ISBN 0387754490.
- Rafael Kiesel, Pietro Totis, and Angelika Kimmig. Efficient knowledge compilation beyond weighted model counting. *Theory and Practice of Logic Programming*, 22(4):505–522, 2022. doi: 10.1017/S147106842200014X.
- Angelika Kimmig, Guy Van den Broeck, and Luc De Raedt. Algebraic model counting. *J. of Applied Logic*, 22(C):46–62, July 2017. ISSN 1570-8683. doi: 10.1016/j.jal.2016.11.031.
- Dieter Kraft. Algorithm 733: TOMP-fortran modules for optimal control calculations. *ACM Transactions on Mathematical Software*, 20(3):262–281, September 1994. ISSN 0098-3500. doi: 10.1145/192115.192124.
- Joohyung Lee and Yi Wang. Weighted rules under the stable model semantics. In Chitta Baral, James P. Delgrande, and Frank Wolter, editors, *Proceedings of the Fifteenth International Conference on Principles of Knowledge Representation and Reasoning*, pages 145–154. AAAI Press, 2016.
- Yuliya Lierler. Constraint answer set programming: Integrational and translational (or smt-based) approaches. *Theory and Practice of Logic Programming*, 23(1):195–225, 2023. doi: 10.1017/S1471068421000478.

- Denis Deratani Mauá and Fábio Gagliardi Cozman. Specifying credal sets with probabilistic answer set programming. In *International Symposium on Imprecise Probabilities and Their Applications*, 2023.
- Aaron Meurer, Christopher P. Smith, Mateusz Paprocki, Ondřej Čertík, Sergey B. Kirpichev, Matthew Rocklin, AMiT Kumar, Sergiu Ivanov, Jason K. Moore, Sartaj Singh, Thilina Rathnayake, Sean Vig, Brian E. Granger, Richard P. Muller, Francesco Bonazzi, Harsh Gupta, Shivam Vats, Fredrik Johansson, Fabian Pedregosa, Matthew J. Curry, Andy R. Terrel, Štěpán Roučka, Ashutosh Saboo, Isuru Fernando, Sumith Kulal, Robert Cimrman, and Anthony Scopatz. SymPy: symbolic computing in python. *PeerJ Computer Science*, 3:e103, January 2017. ISSN 2376-5992. doi: 10.7717/peerj-cs.103.
- Matthias Nickles. Differentiable SAT/ASP. In Elena Bellodi and Tom Schrijvers, editors, *Proceedings of the 5th International Workshop on Probabilistic Logic Programming, PLP 2018, co-located with the 28th International Conference on Inductive Logic Programming (ILP 2018), Ferrara, Italy, September 1, 2018*, volume 2219 of *CEUR Workshop Proceedings*, pages 62–74. CEUR-WS.org, 2018.
- Francesco Pacenza and Jessica Zangari. Extending answer set programming with rational numbers. *arXiv preprint arXiv:2312.04249*, 2023.
- M. J. D. Powell. A direct search optimization method that models the objective and constraint functions by linear interpolation. In Susana Gomez and Jean-Pierre Hennart, editors, *Advances in Optimization and Numerical Analysis*, pages 51–67, Dordrecht, 1994. Springer Netherlands. ISBN 978-94-015-8330-5. doi: 10.1007/978-94-015-8330-5_4.
- Victor Hugo Nascimento Rocha and Fabio Gagliardi Cozman. A credal least undefined stable semantics for probabilistic logic programs and probabilistic argumentation. In Gabriele Kern-Isberner, Gerhard Lakemeyer, and Thomas Meyer, editors, *Proceedings of the 19th International Conference on Principles of Knowledge Representation and Reasoning, KR 2022*, pages 309–319, 8 2022. doi: 10.24963/kr.2022/31.
- Taisuke Sato. A statistical learning method for logic programs with distribution semantics. In Leon Sterling, editor, *ICLP 1995*, pages 715–729. MIT Press, 1995. doi: 10.7551/mitpress/4298.003.0069.
- Leon Sterling and Ehud Shapiro. *The Art of Prolog: Advanced Programming Techniques*. Logic programming. MIT Press, USA, 1994. ISBN 9780262193382.
- Pietro Totis, Luc De Raedt, and Angelika Kimmig. smProbLog: Stable model semantics in problog for probabilistic argumentation. *Theory and Practice of Logic Programming*, pages 1–50, 2023. doi: 10.1017/S147106842300008X.
- Guy Van den Broeck, Ingo Thon, Martijn van Otterlo, and Luc De Raedt. DTProbLog: A decision-theoretic probabilistic Prolog. In Maria Fox and David Poole, editors, *Proceedings of the Twenty-Fourth AAAI Conference on Artificial Intelligence*, pages 1217–1222. AAAI Press, 2010.
- Joost Vennekens, Sofie Verbaeten, and Maurice Bruynooghe. Logic programs with annotated disjunctions. In Bart De-moen and Vladimir Lifschitz, editors, *ICLP 2004*, volume 3131 of *LNCS*, pages 431–445, Berlin, Heidelberg, 2004. Springer. doi: 10.1007/978-3-540-27775-0_30.
- Pauli Virtanen, Ralf Gommers, Travis E. Oliphant, Matt Haberland, Tyler Reddy, David Cournapeau, Evgeni Burovski, Pearu Peterson, Warren Weckesser, Jonathan Bright, Stéfan J. van der Walt, Matthew Brett, Joshua Wilson, K. Jarrod Millman, Nikolay Mayorov, Andrew R. J. Nelson, Eric Jones, Robert Kern, Eric Larson, C J Carey, İlhan Polat, Yu Feng, Eric W. Moore, Jake VanderPlas, Denis Laxalde, Josef Perktold, Robert Cimrman, Ian Henriksen, E. A. Quintero, Charles R. Harris, Anne M. Archibald, Antônio H. Ribeiro, Fabian Pedregosa, Paul van Mulbregt, and SciPy 1.0 Contributors. SciPy 1.0: Fundamental Algorithms for Scientific Computing in Python. *Nature Methods*, 17:261–272, 2020. doi: 10.1038/s41592-019-0686-2.
- Pedro Zuidberg Dos Martires, Anton Dries, and Luc De Raedt. Exact and approximate weighted model integration with probability density functions using knowledge compilation. In *Proceedings of the Thirty-Third AAAI Conference on Artificial Intelligence (AAAI-19)*, pages 7825–7833. AAAI Press, 2019. doi: 10.1609/aaai.v33i01.33017825.